

Amazon_Fine_Food_Reviews_Analysis_KNN

March 7, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to “positive”. Otherwise, it will be set to “negative”.

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from sklearn.neighbors import KNeighborsClassifier
from sklearn.cross_validation import train_test_split
from sklearn.model_selection import TimeSeriesSplit
from sklearn.metrics import roc_auc_score

C:\Users\rites\Anaconda3\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning:
  "This module will be removed in 0.20.", DeprecationWarning)
```

```
In [2]: # using SQLite Table to read data.
```

```

con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000 """, con)

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 100000 """, con)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(-1)
def partition(x):
    if x < 3:
        return 0
    else:
        return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (100000, 10)

```

Out[2]:

```

| | Id | ProductId | UserId | ProfileName | \ |
|---|----|------------|----------------|----------------|------------------|
| 0 | 1 | B001E4KFG0 | A3SGXH7AUHU8GW | delmartian | |
| 1 | 2 | B00813GRG4 | A1D87F6ZCVE5NK | dll pa | |
| 2 | 3 | B000LQOCHO | ABXLMWJIXXAIN | Natalia Corres | "Natalia Corres" |

| | HelpfulnessNumerator | HelpfulnessDenominator | Score | Time | \ |
|---|----------------------|------------------------|-------|------------|---|
| 0 | 1 | 1 | 1 | 1303862400 | |
| 1 | 0 | 0 | 0 | 1346976000 | |
| 2 | 1 | 1 | 1 | 1219017600 | |

| | Summary | Text |
|---|-----------------------|---|
| 0 | Good Quality Dog Food | I have bought several of the Vitality canned d... |
| 1 | Not as Advertised | Product arrived labeled as Jumbo Salted Peanut... |
| 2 | "Delight" says it all | This is a confection that has been around a fe... |

```

In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)

```

```
In [4]: print(display.shape)
        display.head()
```

```
(80668, 7)
```

```
Out [4]:
```

| | UserId | ProductId | ProfileName | Time | Score | \ |
|---|--------------------|------------|------------------------|------------|-------|---|
| 0 | #oc-R115TNMSPFT9I7 | B007Y59HVM | Breyton | 1331510400 | 2 | |
| 1 | #oc-R11D9D7SHXIJB9 | B005HG9ETO | Louis E. Emory "hoppy" | 1342396800 | 5 | |
| 2 | #oc-R11DNU2NBKQ23Z | B007Y59HVM | Kim Cieszykowski | 1348531200 | 1 | |
| 3 | #oc-R1105J5ZVQE25C | B005HG9ETO | Penguin Chick | 1346889600 | 5 | |
| 4 | #oc-R12KPBODL2B5ZD | B007OSBE1U | Christopher P. Presta | 1348617600 | 1 | |

| | Text | COUNT(*) |
|---|---|----------|
| 0 | Overall its just OK when considering the price... | 2 |
| 1 | My wife has recurring extreme muscle spasms, u... | 3 |
| 2 | This coffee is horrible and unfortunately not ... | 2 |
| 3 | This will be the bottle that you grab from the... | 3 |
| 4 | I didnt like this coffee. Instead of telling y... | 2 |

```
In [5]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out [5]:
```

| | UserId | ProductId | ProfileName | Time | \ |
|-------|---------------|------------|---------------------------------|------------|---|
| 80638 | AZY10LLTJ71NX | B006P7E5ZI | undertheshrine "undertheshrine" | 1334707200 | |

| | Score | Text | COUNT(*) |
|-------|-------|---|----------|
| 80638 | 5 | I was recommended to try green tea extract to ... | 5 |

```
In [6]: display['COUNT(*)'].sum()
```

```
Out [6]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [7]: display= pd.read_sql_query("""
        SELECT *
        FROM Reviews
        WHERE Score != 3 AND UserId="AR5J8UI46CURR"
        ORDER BY ProductID
        """, con)
        display.head()
```

```

Out [7]:
      Id  ProductId  UserId  ProfileName  HelpfulnessNumerator  \
0   78445  B000HDL1RQ  AR5J8UI46CURR  Geetha Krishnan             2
1  138317  B000HDOPYC  AR5J8UI46CURR  Geetha Krishnan             2
2  138277  B000HDOPYM  AR5J8UI46CURR  Geetha Krishnan             2
3   73791  B000HDOPZG  AR5J8UI46CURR  Geetha Krishnan             2
4  155049  B000PAQ75C  AR5J8UI46CURR  Geetha Krishnan             2

      HelpfulnessDenominator  Score  Time  \
0                2          5  1199577600
1                2          5  1199577600
2                2          5  1199577600
3                2          5  1199577600
4                2          5  1199577600

      Summary  \
0  LOACKER QUADRATINI VANILLA WAFERS
1  LOACKER QUADRATINI VANILLA WAFERS
2  LOACKER QUADRATINI VANILLA WAFERS
3  LOACKER QUADRATINI VANILLA WAFERS
4  LOACKER QUADRATINI VANILLA WAFERS

      Text
0  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4  DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

```

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```

In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)

In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape

```

```
Out[9]: (87775, 10)
```

```
In [10]: #Checking to see how much % of data still remains
         (final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[10]: 87.775
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [11]: display= pd.read_sql_query("""
        SELECT *
        FROM Reviews
        WHERE Score != 3 AND Id=44737 OR Id=64422
        ORDER BY ProductID
        """, con)
```

```
display.head()
```

```
Out[11]:
```

| | Id | ProductId | UserId | ProfileName | \ |
|---|-------|------------|----------------|----------------|----------|
| 0 | 64422 | B000MIDR0Q | A161DK06JJMCYF | J. E. Stephens | "Jeanne" |
| 1 | 44737 | B001EQ55RW | A2V0I904FH7ABY | | Ram |

| | HelpfulnessNumerator | HelpfulnessDenominator | Score | Time | \ |
|---|----------------------|------------------------|-------|------|------------|
| 0 | | 3 | 1 | 5 | 1224892800 |
| 1 | | 3 | 2 | 4 | 1212883200 |

| | Summary | \ |
|---|--|---|
| 0 | Bought This for My Son at College | |
| 1 | Pure cocoa taste with crunchy almonds inside | |

| | Text |
|---|---|
| 0 | My son loves spaghetti so I didn't hesitate or... |
| 1 | It was almost a 'love at first bite' - the per... |

```
In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [13]: #Before starting the next phase of preprocessing lets see the number of entries left
         print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

```
(87773, 10)
```

```
Out[13]: 1    73592
         0    14181
         Name: Score, dtype: int64
```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [14]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its

=====

The Candy Blocks were a nice visual for the Lego Birthday party but the candy has little taste

=====

was way to hot for my blood, took a bite and did a jig lol

=====

My dog LOVES these treats. They tend to have a very strong fish oil smell. So if you are afraid

=====

```
In [15]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
```

```

sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)

```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its

```

In [16]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
sent_0 = soup.get_text()
print(sent_0)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
sent_1000 = soup.get_text()
print(sent_1000)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
sent_1500 = soup.get_text()
print(sent_1500)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
sent_4900 = soup.get_text()
print(sent_4900)

```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its

=====

The Candy Blocks were a nice visual for the Lego Birthday party but the candy has little taste

=====

was way to hot for my blood, took a bite and did a jig lol

=====

My dog LOVES these treats. They tend to have a very strong fish oil smell. So if you are afraid

```

In [17]: # https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)
    phrase = re.sub(r"wont", "will not", phrase) # in some words apostrophe is missing
    phrase = re.sub(r"its", "it is", phrase)
    phrase = re.sub(r"Its", "It is", phrase)

```



```
'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'a
'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'to
's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", '
've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't
"hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mi
"mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
'won', "won't", 'wouldn', "wouldn't"])
```

```
In [22]: # Combining all the above students
import itertools
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)

    #https://www.analyticsvidhya.com/blog/2014/11/text-data-cleaning-steps-python/
    # This removes words such as aawwww or happpyyy or awsoooooomee etc
    sentence = ''.join(''.join(s)[:2] for _, s in itertools.groupby(sentence))

    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```

100%|| 87773/87773 [01:07<00:00, 1301.29it/s]

```
In [23]: smpl = " aaaaww aaaww aaa aawwww"
sent = ''.join(''.join(s)[:2] for _, s in itertools.groupby(smpl))
print(sent)
```

aaww aaww aa aaww

```
In [24]: preprocessed_reviews[4900]
```

```
Out[24]: 'dog loves treats tend strong fish oil smell afraid fishy smell not get think dog like'
```

[3.2] Preprocessing Review Summary

```
In [25]: ## Similarly performing preprocessing for review summary also.
```

```
preprocessed_summary=[]

for sent in tqdm(final['Summary'].values):
```

```

sent = re.sub(r"http\S+", "", sent)
sent = BeautifulSoup(sent, 'lxml').get_text()
sent = decontracted(sent)
sent = re.sub(r"\S+\d\S+", "", sent).strip()
sent = re.sub(r"[^A-Za-z0-9]+", "", sent)

#https://www.analyticsvidhya.com/blog/2014/11/text-data-cleaning-steps-python/
# This removes words such as aawwww or happpyyy or awsoooooomee etc
sent = ' '.join(' '.join(s)[:2] for _, s in itertools.groupby(sent))

# https://gist.github.com/sebleier/554280
sent = ' '.join(w.lower() for w in sent.split() if w.lower() not in stopwords)
preprocessed_summary.append(sent.strip())

```

100%|| 87773/87773 [00:34<00:00, 2512.46it/s]

In [26]: preprocessed_summary[4900]

Out[26]: 'great value'

In [27]: final['Summary'].values[4900]

Out[27]: 'Great value'

```

In [28]: # Removing those words which are of lenght 2
# This will remove non relevant words then we will perform featurization
cleaned_reviews = []
for sent in preprocessed_reviews:
    sentence = ' '.join(w for w in sent.split() if len(w)>2)
    cleaned_reviews.append(sentence.strip())

```

In [29]: print(cleaned_reviews[0])

dogs loves chicken product china not buying anymore hard find chicken products made usa one no

```

In [30]: final["Cleaned_review"] = cleaned_reviews
final.head(5)

```

```

Out[30]:
      Id  ProductId  UserId  ProfileName \
22620  24750  2734888454  A13ISQV0U9GZIC  Sandikaye
22621  24751  2734888454  A1C298ITT645B6  Hugh G. Pritchard
70677  76870  B00002N8SM  A19Q006CSFT011  Arlielle
70676  76869  B00002N8SM  A1FYH4S02BW7FN  wonderer
70675  76868  B00002N8SM  AUE8TB5VHS6ZV  eyeofthestorm

      HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
22620                    1                      1      0  1192060800

```

| | | | | |
|-------|---|---|---|------------|
| 22621 | 0 | 0 | 1 | 1195948800 |
| 70677 | 0 | 0 | 0 | 1288396800 |
| 70676 | 0 | 0 | 0 | 1290038400 |
| 70675 | 0 | 0 | 0 | 1306972800 |

| | Summary \ |
|-------|---|
| 22620 | made in china |
| 22621 | Dog Lover Delites |
| 70677 | only one fruitfly stuck |
| 70676 | Doesn't work!! Don't waste your money!! |
| 70675 | A big rip off |

| | Text \ |
|-------|---|
| 22620 | My dogs loves this chicken but its a product f... |
| 22621 | Our dogs just love them. I saw them in a pet ... |
| 70677 | I had an infestation of fruitflies, they were ... |
| 70676 | Worst product I have gotten in long time. Woul... |
| 70675 | I wish I'd read the reviews before making this... |

| | Cleaned_review |
|-------|---|
| 22620 | dogs loves chicken product china not buying an... |
| 22621 | dogs love saw pet store tag attached regarding... |
| 70677 | infestation fruitflies literally everywhere fl... |
| 70676 | worst product gotten long time would rate star... |
| 70675 | wish would read reviews making purchase basica... |

5 [4] Featurization

5.1 [4.1] BAG OF WORDS

In [36]: *#BoW*

```
count_vect1 = CountVectorizer() #in scikit-learn
count_vect1.fit(cleaned_reviews)
print("some feature names ", count_vect1.get_feature_names()[:10])
print('='*50)

final_counts = count_vect1.transform(cleaned_reviews)
print("the type of count vectorizer ",type(final_counts))
print("the shape of out text BOW vectorizer ",final_counts.get_shape())
print("the number of unique words ", final_counts.get_shape()[1])
```

```
some feature names  ['aaa', 'aaah', 'aaahh', 'aaaww', 'aachen', 'aadp', 'aaf', 'aafco', 'aah',
=====
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer  (87773, 54095)
the number of unique words  54095
```

In [37]: *#bi-gram, tri-gram and n-gram*

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (87773, 112780)
the number of unique words including both unigrams and bigrams 112780
['aafco', 'aback', 'abandon', 'abandoned', 'abc', 'abdomen', 'abdominal', 'abdominal pain', 'al
```

```
In [38]: # tfidf on unigrams
```

```
some sample features(unique words in the corpus) ['aaa', 'aaah', 'aaahh', 'aaaww', 'aachen', 'aachn']
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (87773, 54095)
the number of unique words including both unigrams and bigrams 54095
```

```
tf_idf_vect2 = TfidfVectorizer(ngram_range=(1,2),min_df=3)
tf_idf_vect2.fit(cleaned_reviews)
print("some sample features(unique words in the corpus)",tf_idf_vect2.get_feature_names())
print('='*50)
```

```

final_tf_idf2 = tf_idf_vect2.transform(cleaned_reviews)
print("the type of count vectorizer ",type(final_tf_idf2))
print("the shape of out text TFIDF vectorizer ",final_tf_idf2.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf2)

some sample features(unique words in the corpus) ['aafco', 'aafco dog', 'aah', 'aahs', 'aback']
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (87773, 213055)
the number of unique words including both unigrams and bigrams 213055

```

5.4 [4.4] Word2Vec

In [40]: *# Train your own Word2Vec model using your own text corpus*

```

i=0
list_of_sentence=[]
for sentence in tqdm(cleaned_reviews):
    list_of_sentence.append(sentence.split())

```

100%|| 87773/87773 [00:00<00:00, 107758.27it/s]

In [32]: `outfile = open("list_of_sentence","wb")`
`pickle.dump(list_of_sentence,outfile)`
`outfile.close()`

In [41]: *# Using Google News Word2Vectors*

```

# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.

# http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzZPY
# you can comment this whole cell
# or change these variable according to your need

is_your_ram_gt_16g=False
want_to_use_google_w2v =False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occurred atleast 5 times
    w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))

```

```

print('='*50)
print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.b
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have gogole's word2vec file, keep want_to_train_w2v = True, t

[('fantastic', 0.8564410209655762), ('good', 0.8339499235153198), ('terrific', 0.8293559551239
=====
[('greatest', 0.8057908415794373), ('tastiest', 0.7508324384689331), ('best', 0.71000838279724

In [42]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occured minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])

```

number of words that occured minimum 5 times 17061
sample words ['dogs', 'loves', 'chicken', 'product', 'china', 'not', 'buying', 'anymore', 'ha

5.5 [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

```

In [52]: # average Word2Vec
# compute average word2vec for each review.
sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in list_of_sentence: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need t
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))

```

KeyboardInterrupt

Traceback (most recent call last)

```

<ipython-input-52-6e7e3cb11537> in <module>()
      8         if word in w2v_words:
      9             vec = w2v_model.wv[word]
----> 10             sent_vec += vec
      11             cnt_words += 1
      12         if cnt_words != 0:

```

KeyboardInterrupt:

[4.4.1.2] TFIDF weighted W2v

```

In [43]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
tfidf_matrix = model.fit_transform(cleaned_reviews)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

In [44]: # TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this l
row=0;
for sent in list_of_sentence: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf values of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += 1
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1

```

KeyboardInterrupt

Traceback (most recent call last)


```

<ipython-input-44-acbebc07d9a9> in <module>()
    16             # sent.count(word) = tf values of word in this review
    17             tf_idf = dictionary[word]*(sent.count(word)/len(sent))
---> 18             sent_vec += (vec * tf_idf)
    19             weight_sum += 1
    20             if weight_sum != 0:

```

KeyboardInterrupt:

```

In [45]: print(len(tfidf_sent_vectors))
         print(len(tfidf_sent_vectors[0]))

```

```

64957
50

```

```

In [46]: # Function to plot confusion matrix
def confusion_matrix_plot(test_y, predict_y):
    # C stores the confusion matrix
    C = confusion_matrix(test_y, predict_y)

    # Class labels
    labels_x = ["Predicted No", "Predicted Yes"]
    labels_y = ["Original No", "Original Yes"]

    cmap=sns.light_palette("orange")
    print("Confusion matrix")
    plt.figure(figsize=(4,4))
    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels_x, yticklabels=labels_y)
    plt.show()

```

```

In [47]: # Function to plot roc curve

```

```

def plot_roc_curve(Y_test, predict_y_test, Y_train, predict_y_train):
    fpr1, tpr1, threshold1 = roc_curve(Y_test, predict_y_test) # For test dataset
    fpr2, tpr2, threshold2 = roc_curve(Y_train, predict_y_train) # For train dataset

    plt.plot([0,1],[0,1])
    plt.plot(fpr1, tpr1, label="Validation AUC")
    plt.plot(fpr2, tpr2, label="Train AUC")
    plt.xlabel("fpr")
    plt.ylabel("tpr")
    plt.legend()
    plt.show()

```

```

In [48]: # Plotting graph of auc and parameter for training and cross validation error
param = [1,3,5,7,9,11,13,15,17,19,21,23,25,27,29]

```

```
def plot_knn_vs_auc(train_auc_list,cv_auc_list):
    plt.plot(param,train_auc_list,label="Train AUC")
    plt.xlabel("Parameter for K-NN")
    plt.ylabel("Area Under Curve")
    plt.plot(param,cv_auc_list,label="Validation AUC")
    plt.legend()
    plt.show()
```

6 [5] Assignment 3: KNN

Apply Knn(brute force version) on these feature sets

SET 1:Review text, preprocessed one converted into vectors using (BOW)

SET 2:Review text, preprocessed one converted into vectors using (TFIDF)

SET 3:Review text, preprocessed one converted into vectors using (AVG W2v)

SET 4:Review text, preprocessed one converted into vectors using (TFIDF W2v)

Apply Knn(kd tree version) on these feature sets NOTE: sklearn implementation of kd-tree accepts only dense matrices, you need to convert the sparse matrices of CountVectorizer/TfidfVectorizer into dense matrices. You can convert sparse matrices to dense using .toarray() attribute. For more information please visit this link

SET 5:Review text, preprocessed one converted into vectors using (BOW) but with restriction on maximum features generated.

```
</li>
<li><font color='red'>SET 6:</font>Review text, preprocessed one converted into vectors
<pre>
    tf_idf_vect = TfidfVectorizer(min_df=10, max_features=500)
    tf_idf_vect.fit(preprocessed_reviews)
</pre>
</li>
<li><font color='red'>SET 3:</font>Review text, preprocessed one converted into vectors
<li><font color='red'>SET 4:</font>Review text, preprocessed one converted into vectors
</ul>
</li>
<br>
<li><strong>The hyper paramter tuning(find best K)</strong>
    <ul>
<li>Find the best hyper parameter which will give the maximum <a href='https://www.appliedaicon
<li>Find the best hyper paramter using k-fold cross validation or simple cross validation data
<li>Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this ta
    </ul>
</li>
<br>
<li>
<strong>Representation of results</strong>
    <ul>
<li>You need to plot the performance of model both on train data and cross validation data for
```

```

<img src='train_cv_auc.JPG' width=300px></li>
<li>Once after you found the best hyper parameter, you need to train your model with it, and f
<img src='train_test_auc.JPG' width=300px></li>
<li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicourse.
<img src='confusion_matrix.png' width=300px></li>
    </ul>
</li>
<br>
<li><strong>Conclusion</strong>
    <ul>
<li>You need to summarize the results at the end of the notebook, summarize it in the table fo
    <img src='summary.JPG' width=400px>
</li>
    </ul>

```

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this link.

6.1 [5.1] Applying KNN brute force

6.1.1 [5.1.1] Applying KNN brute force on BOW, SET 1

```

In [49]: from sklearn.cross_validation import train_test_split

        # cleaned_reviews contains all the required reviews
        # Splitting cleaned_reviews into train and test dataset

        X = cleaned_reviews
        Y = final['Score']

        X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.3,random_state=42)
        print(len(X_train),len(Y_train),len(X_test),len(Y_test))

61441 61441 26332 26332

In [50]: # Now we will vectorize train and test datasets separately using BagofWords
        # Use fit_transform to vectorize train dataset and transform to vectorize test dataset
        count_vect2 = CountVectorizer(max_features=2000)
        X_train = count_vect2.fit_transform(X_train)
        X_test = count_vect2.transform(X_test)
        print(X_train.shape,X_test.shape)

```

(61441, 2000) (26332, 2000)

```
In [52]: # In this section we will calculate training error.
# To calculate training error you have to train model using training data and
# Then test the same model on training data and compare the predicted labels with actual
from sklearn.model_selection import TimeSeriesSplit
from sklearn.metrics import roc_auc_score
from sklearn.neighbors import KNeighborsClassifier

param_list = [1,3,5,7,9,11,13,15,17,19,21,23,25,27,29]
train_auc_list1 = [] # This contains area under curve value of first batch corresponding to param_list
train_auc_list2 = [] # for second batch
train_auc_list3 = [] # for third batch

# Testing whole training data at once takes a lot of memory which takes a lot of time
# There fore we are dividing training data into 3 parts and then we will calculate AUROC

x_train_1 = X_train[0:20000][:] # Row 0 to 19999 and all columns
x_train_2 = X_train[20000:40000][:]
x_train_3 = X_train[40000:61441][:]

y_train_1 = Y_train[0:20000][:] # Row 0 to 19999 and all columns
y_train_2 = Y_train[20000:40000][:]
y_train_3 = Y_train[40000:61441][:]

# Calculating training error for first batch
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="brute",leaf_size=30)
    clf.fit(x_train_1,y_train_1)

    pre_probab = clf.predict_proba(x_train_1)[:,-1] # Returns probability of positive class

    auc = roc_auc_score(y_train_1,pre_probab)
    train_auc_list1.append(auc)

# Calculating training error for second batch
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="brute",leaf_size=30)
    clf.fit(x_train_2,y_train_2)

    pre_probab = clf.predict_proba(x_train_2)[:,-1]

    auc = roc_auc_score(y_train_2,pre_probab)
    train_auc_list2.append(auc)

# Calculating training error for third batch
```

```

for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="brute",leaf_size=30)
    clf.fit(x_train_3,y_train_3)

    pre_probab = clf.predict_proba(x_train_3)[:,-1]

    auc = roc_auc_score(y_train_3,pre_probab)
    train_auc_list3.append(auc)

100%|| 15/15 [10:10<00:00, 35.93s/it]
100%|| 15/15 [06:07<00:00, 24.80s/it]
100%|| 15/15 [06:35<00:00, 25.38s/it]

In [53]: # Combining training result of each batch together
train_auc_list = [(x+y+z)/3 for x,y,z in zip(train_auc_list1,train_auc_list2,train_auc_list3)]

In [54]: # We will do time based splitting and do 10 fold cross validation
# This is done as reviews keeps changing with time and hence time based splitting is used

# Time series object
tscv = TimeSeriesSplit(n_splits=10)

cv_auc_list = [] # will contain cross validation AUC corresponding to each k

for k in range(1,30,2):
    # KNN Classifier
    clf = KNeighborsClassifier(n_neighbors=k,algorithm='brute',leaf_size=30)
    i=0
    auc=0.0
    for train_index,test_index in tscv.split(X_train):
        x_train = X_train[0:train_index[-1]][:] # row 0 to train_index(excluding)
        y_train = Y_train[0:train_index[-1]][:] # row 0 to train_index(excluding)
        x_test = X_train[train_index[-1]:test_index[-1]][:] # row from train_index to test_index(excluding)
        y_test = Y_train[train_index[-1]:test_index[-1]][:] # row from train_index to test_index(excluding)

        clf.fit(x_train,y_train)


        predict_probab = clf.predict_proba(x_test)[:,-1]
        i += 1
        auc += roc_auc_score(y_test,predict_probab)

    cv_auc_list.append(auc/i) # Storing AUC value

In [55]: import matplotlib.pyplot as plt

```

```
# Plotting graph of auc and parameter for training and cross validation error
plot_knn_vs_auc(train_auc_list,cv_auc_list)
```



Observing the graph we will select a k for which AUC is not very high in training error plot to avoid overfitting and select a k for which AUC is not very low in Cross-validation error to avoid underfitting. Therefore we are selecting k = 25 .

```
In [56]: # Training final model on best auc and taking k = 25
```

```
# Training one model with all the data hangs the PC .
# Therefore we will divide data into 3 parts and then train three separate models.
final_clf1 = KNeighborsClassifier(n_neighbors=25,algorithm='brute',leaf_size=30)
final_clf1.fit(x_train_1,y_train_1)
predict_probab_1 = final_clf1.predict_proba(X_test)[: ,1] # This returns only probabil
predict_y1 = final_clf1.predict(X_test)
predict_y_train1 = final_clf1.predict(x_train_1)

final_clf2 = KNeighborsClassifier(n_neighbors=25,algorithm='brute',leaf_size=30)
final_clf2.fit(x_train_2,y_train_2)
predict_probab_2 = final_clf2.predict_proba(X_test)[: ,1] # This returns only probabil
predict_y2 = final_clf2.predict(X_test)
predict_y_train2 = final_clf2.predict(x_train_2)

final_clf3 = KNeighborsClassifier(n_neighbors=25,algorithm='brute',leaf_size=30)
final_clf3.fit(x_train_3,y_train_3)
predict_probab_3 = final_clf3.predict_proba(X_test)[: ,1] # This returns only probabil
```

```

predict_y3 = final_clf3.predict(X_test)
predict_y_train3 = final_clf3.predict(x_train_3)

# Now merging, n_jobs=3 all the three probability scores into one

predict_probab = [(x+y+z)/3 for x,y,z in zip(predict_probab_1,predict_probab_2,predict_probab_3)]

auc = roc_auc_score(Y_test,predict_probab)
print("Final AUC is ::{:.2f}".format(auc))

```

Final AUC is ::0.73

```


In [57]: # combining labels predicted by all the models
predict_y = [1 if (x+y+z)>= 2 else 0 for x,y,z in zip(predict_y1,predict_y2,predict_y3)]

In [58]: # Combining results of train dataset evaluation
predict_y_train = np.concatenate((predict_y_train1,predict_y_train2),axis=None)
predict_y_train = np.concatenate((predict_y_train,predict_y_train3),axis=None)


In [59]: # Plotting confusion matrix
confusion_matrix_plot(Y_test,predict_y)

```

Confusion matrix



```
In [60]: # Plotting ROC Curve
plot_roc_curve(Y_test,predict_y,Y_train,predict_y_train)
```



6.1.2 [5.1.2] Applying KNN brute force on TFIDF

```
In [99]: # In this section Tfidf will be used for vectorization
# Splitting datasets into train and test datasets
```

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.3,random_state=42)
```

```
# Initializinf TFidf
```

```
tf_idf_vect2 = TfidfVectorizer(max_features=2000)
```

```
# Now we will vectorize train and test datasets separately using Tfidf
```

```
# Use fit_transform to vectorize train dataset and transform to vectorize test dataset
```

```
X_train = tf_idf_vect2.fit_transform(X_train)
```

```
X_test = tf_idf_vect2.transform(X_test)
```

```
In [100]: param_list = [1,3,5,7,9,11,13,15,17,19,21,23,25,27,29]
```

```
train_auc_list1 = [] # This contains area under curve value of first batch correspond
```

```
train_auc_list2 = [] # for second batch
```

```
train_auc_list3 = [] # for third batch
```

```
# Testing whole training data at once takes a lot of memory which takes a lot of time
```

```
# Therefore we are dividing training data into 3 parts and then we will calculate A
```



```

x_train_1 = X_train[0:20000][:] # Row 0 to 19999 and all columns
x_train_2 = X_train[20000:40000][:]
x_train_3 = X_train[40000:61441][:]

y_train_1 = Y_train[0:20000][:] # Row 0 to 19999 and all columns
y_train_2 = Y_train[20000:40000][:]
y_train_3 = Y_train[40000:61441][:]

# Calculating training error for first batch
for k in range(1,30,2):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="brute",leaf_size=30)
    clf.fit(x_train_1,y_train_1)

    pre_probab = clf.predict_proba(x_train_1)[: ,1] # Returns probability of positive

    auc = roc_auc_score(y_train_1,pre_probab)
    train_auc_list1.append(auc)

# Calculating training error for second batch
for k in range(1,30,2):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="brute",leaf_size=30)
    clf.fit(x_train_2,y_train_2)

    pre_probab = clf.predict_proba(x_train_2)[: ,1]

    auc = roc_auc_score(y_train_2,pre_probab)
    train_auc_list2.append(auc)

# Calculating training error for third batch
for k in range(1,30,2):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="brute",leaf_size=30)
    clf.fit(x_train_3,y_train_3)

    pre_probab = clf.predict_proba(x_train_3)[: ,1]

    auc = roc_auc_score(y_train_3,pre_probab)
    train_auc_list3.append(auc)

# Combining training result of each batch together
train_auc_list = [(x+y+z)/3 for x,y,z in zip(train_auc_list1,train_auc_list2,train_auc_list3)]

```

In [101]: # Performing time series split cross validation

```

auc_list=[]

for k in range(1,30,2):
    # KNN Classifier

```

```

clf = KNeighborsClassifier(n_neighbors=k,algorithm='brute',leaf_size=30,n_jobs=3)
i=0
auc=0.0
for train_index,test_index in tscv.split(X_train):
    x_train = X_train[0:train_index[-1]][:] # row 0 to train_index(excluding)
    y_train = Y_train[0:train_index[-1]][:] # row 0 to train_index(excluding)
    x_test = X_train[train_index[-1]:test_index[-1]][:] # row from train_index to
    y_test = Y_train[train_index[-1]:test_index[-1]][:] # row from train_index to

    clf.fit(x_train,y_train)

    predict_probab = clf.predict_proba(x_test)[:,-1]
    i += 1
    auc += roc_auc_score(y_test,predict_probab)

auc_list.append(auc/i)


```

In [102]: `import matplotlib.pyplot as plt`

```

# Plotting graph of auc and parameter for training and cross validation error
plot_knn_vs_auc(train_auc_list,auc_list)

```



```

In [103]: # Training final model on best auc and taking k = 21

# Training one model with all the data hangs the PC .
# Therefore we will divide data into 3 parts and then train three separate models.
final_clf1 = KNeighborsClassifier(n_neighbors=21,algorithm='brute',leaf_size=30)
final_clf1.fit(x_train_1,y_train_1)
predict_probab_1 = final_clf1.predict_proba(X_test)[:,:1] # This returns only probability
predict_y1 = final_clf1.predict(X_test)
predict_y_train1 = final_clf1.predict(x_train_1)

final_clf2 = KNeighborsClassifier(n_neighbors=21,algorithm='brute',leaf_size=30)
final_clf2.fit(x_train_2,y_train_2)
predict_probab_2 = final_clf2.predict_proba(X_test)[:,:1] # This returns only probability
predict_y2 = final_clf2.predict(X_test)
predict_y_train2 = final_clf2.predict(x_train_2)

final_clf3 = KNeighborsClassifier(n_neighbors=21,algorithm='brute',leaf_size=30)
final_clf3.fit(x_train_3,y_train_3)
predict_probab_3 = final_clf3.predict_proba(X_test)[:,:1] # This returns only probability
predict_y3 = final_clf3.predict(X_test)
predict_y_train3 = final_clf3.predict(x_train_3)

# Now merging all the three probability scores into one
predict_probab = [(x+y+z)/3 for x,y,z in zip(predict_probab_1,predict_probab_2,predict_probab_3)]
predict_y = [1 if(x+y+z>= 2) else 0 for x,y,z in zip(predict_y1,predict_y2,predict_y3)]

# Combining results of train dataset evaluation
predict_y_train = np.concatenate((predict_y_train1,predict_y_train2),axis=None)
predict_y_train = np.concatenate((predict_y_train,predict_y_train3),axis=None)

auc = roc_auc_score(Y_test,predict_probab)
print("Final AUC is ::{:0.2f}".format(auc))

```


Final AUC is ::0.51

```


In [104]: # plotting confusion matrix
confusion_matrix_plot(Y_test,predict_y)

```

Confusion matrix



```
In [105]: # Plotting roc curve  
plot_roc_curve(Y_test,predict_y,Y_train,predict_y_train)
```



6.1.3 [5.1.3] Applying KNN brute force on AVG W2V, SET 3

```
In [106]: # In this section avg_w2v will be used for vectorization
# Splitting datasets into train and test datasets
Y = final['Score']
X_train,X_test,Y_train,Y_test = train_test_split(list_of_sentence,Y,test_size=0.3,ran
print(X_train[0])

['use', 'beans', 'espresso', 'machine', 'love', 'taste', 'straight', 'espresso', 'coffee', 'fi
```

```
In [107]: # Now we will vectorize train dataset usin avg_w2v
train_sent_vectors = []; # the avg-w2v for each sentence/review is stored in this li
for sent in X_train: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    train_sent_vectors.append(sent_vec)
print(len(train_sent_vectors))
print(len(train_sent_vectors[0]))
```

```
61441
50
```

```
In [108]: # Vectorization of test dataset using avg_w2v

test_sent_vectors = []; # the avg-w2v for each sentence/review is stored in this lis
for sent in X_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    test_sent_vectors.append(sent_vec)
print(len(test_sent_vectors))
print(len(test_sent_vectors[0]))
```

```
26332
50
```

```

In [109]: param_list = [1,3,5,7,9,11,13,15,17,19,21,23,25,27,29]
          train_auc_list1 = []    # This contains area under curve value of first batch corresp
          train_auc_list2 = []    # for second batch
          train_auc_list3 = []    # for third batch

          # Testing whole training data at once takes a lot of memory which takes a lot of time
          # There fore we are dividing training data into 3 parts and then we will calculate A

          x_train_1 = train_sent_vectors[0:20000][:] # Row 0 to 19999 and all columns
          x_train_2 = train_sent_vectors[20000:40000][:]
          x_train_3 = train_sent_vectors[40000:61441][:]

          y_train_1 = Y_train[0:20000][:] # Row 0 to 19999 and all columns
          y_train_2 = Y_train[20000:40000][:]
          y_train_3 = Y_train[40000:61441][:]

          # Calculating training error for first batch
          for k in tqdm(range(1,30,2)):
              clf = KNeighborsClassifier(n_neighbors=k,algorithm="brute")
              clf.fit(x_train_1,y_train_1)

              pre_probab = clf.predict_proba(x_train_1)[:,:1] # Returns probability of positive

              auc = roc_auc_score(y_train_1,pre_probab)
              train_auc_list1.append(auc)

          # Calculating training error for second batch
          for k in tqdm(range(1,30,2)):
              clf = KNeighborsClassifier(n_neighbors=k,algorithm="brute")
              clf.fit(x_train_2,y_train_2)

              pre_probab = clf.predict_proba(x_train_2)[:,:1]

              auc = roc_auc_score(y_train_2,pre_probab)
              train_auc_list2.append(auc)

          # Calculating training error for third batch
          for k in tqdm(range(1,30,2)):
              clf = KNeighborsClassifier(n_neighbors=k,algorithm="brute")
              clf.fit(x_train_3,y_train_3)

              pre_probab = clf.predict_proba(x_train_3)[:,:1]

              auc = roc_auc_score(y_train_3,pre_probab)
              train_auc_list3.append(auc)

```

100%|| 15/15 [02:54<00:00, 12.03s/it]

```
100%|| 15/15 [02:50<00:00, 11.73s/it]
100%|| 15/15 [03:19<00:00, 14.13s/it]
```

```
In [110]: # Combining training result of each batch together
```

```
train_auc_list = [(x+y+z)/3 for x,y,z in zip(train_auc_list1,train_auc_list2,train_auc_list3)]
```

```
In [111]: # 10 fold cross validation using time series splitting
```

```
from sklearn.model_selection import TimeSeriesSplit
```

```
tscv = TimeSeriesSplit(n_splits=10)
```

```
auc_list=[]
```

```
for k in range(1,30,2):
```

```
    # KNN Classifier
```

```
    clf = KNeighborsClassifier(n_neighbors=k,algorithm='brute',leaf_size=30)
```

```
    i=0
```

```
    auc=0.0
```

```
    for train_index,test_index in tscv.split(train_sent_vectors):
```

```
        x_train = train_sent_vectors[0:train_index[-1]][:] # row 0 to train_index(excluding)
```

```
        y_train = Y_train[0:train_index[-1]][:] # row 0 to train_index(excluding)
```

```
        x_test = train_sent_vectors[train_index[-1]:test_index[-1]][:] # row from train_index to test_index(excluding)
```

```
        y_test = Y_train[train_index[-1]:test_index[-1]][:] # row from train_index to test_index(excluding)
```

```
        clf.fit(x_train,y_train)
```

```
        predict_probab = clf.predict_proba(x_test)[:,-1]
```


```
        i += 1
```

```
        auc += roc_auc_score(y_test,predict_probab)
```

```
    auc_list.append(auc/i)
```

```
In [113]: # Plotting graph of auc and parameter for training and cross validation error
```

```
plot_knn_vs_auc(train_auc_list, auc_list)
```



```
In [115]: # Training final model on best auc and taking k = 25
```

```
# Training one model with all the data hangs the PC .
```

```
# Therefore we will divide data into 3 parts and then train three seperate models.
```

```
final_clf1 = KNeighborsClassifier(n_neighbors=30,algorithm='brute',leaf_size=40)
```

```
final_clf1.fit(x_train_1,y_train_1)
```

```
predict_probab_1 = final_clf1.predict_proba(test_sent_vectors)[: ,1] # This returns o
```

```
predict_y1 = final_clf1.predict(test_sent_vectors)
```

```
predict_y_train1 = final_clf1.predict(x_train_1)
```

```
final_clf2 = KNeighborsClassifier(n_neighbors=30,algorithm='brute',leaf_size=40)
```

```
final_clf2.fit(x_train_2,y_train_2)
```

```
predict_probab_2 = final_clf2.predict_proba(test_sent_vectors)[: ,1] # This returns o
```

```
predict_y2 = final_clf2.predict(test_sent_vectors)
```

```
predict_y_train2 = final_clf2.predict(x_train_2)
```

```
final_clf3 = KNeighborsClassifier(n_neighbors=30,algorithm='brute',leaf_size=40)
```

```
final_clf3.fit(x_train_3,y_train_3)
```

```
predict_probab_3 = final_clf3.predict_proba(test_sent_vectors)[: ,1] # This returns o
```

```
predict_y3 = final_clf3.predict(test_sent_vectors)
```

```
predict_y_train3 = final_clf3.predict(x_train_3)
```

```
# Now merging all the three probability scores into one
```

```
predict_probab = [(x+y+z)/3 for x,y,z in zip(predict_probab_1,predict_probab_2,predict_probab_3)]
```



```

predict_y = [1 if(x+y+z>= 2) else 0 for x,y,z in zip(predict_y1,predict_y2,predict_y3)]

auc = roc_auc_score(Y_test,predict_probab)
print("Final AUC is ::{:0.2f}".format(auc))

```

Final AUC is ::0.90

```

In [116]: predict_y_train = [] # will store predicted class labels of combined predicted labels

# Combining predicted labels of train dataset evaluation
# Appending predicted labels of first model
for i in predict_y_train1:
    predict_y_train.append(i)

# Appending predicted labels of second model
for i in predict_y_train2:
    predict_y_train.append(i)

# Appending predicted labels of third model
for i in predict_y_train3:
    predict_y_train.append(i)

print(len(predict_y_train))

```


61441

```


In [117]: # Plotting confusion matrix
confusion_matrix_plot(Y_test,predict_y)

```

Confusion matrix



```
In [118]: # Plotting roc  
plot_roc_curve(Y_test,predict_y,Y_train,predict_y_train)
```



6.1.4 [5.1.4] Applying KNN brute force on TFIDF W2V

```
In [103]: Y = final['Score'] # Contains labels of data points
          X_train,X_test,Y_train,Y_test = train_test_split(cleaned_reviews,Y,test_size=0.3,ran

In [104]: # Vectorizing train dataset using tfidf
          model = TfidfVectorizer()
          model.fit_transform(X_train)
          tfidf_feat1 = model.get_feature_names()

          # This will map word with their tfidf only for train dataset
          dictionary1 = dict(zip(model.get_feature_names(), list(model.idf_)))

In [105]: # Vectorizing test dataset using tfidf
          model.transform(X_test)
          tfidf_feat2 = model.get_feature_names()

          # This will map word with their tfidf only for test dataset
          dictionary2 = dict(zip(model.get_feature_names(), list(model.idf_)))

In [106]: # Vectorizing train dataset

          # TF-IDF weighted Word2Vec
          # tfidf words/col-names
          # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

          train_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in

          for sent in X_train: # for each review/sentence
              sent_vec = np.zeros(50) # as word vectors are of zero length
              weight_sum = 0; # num of words with a valid vector in the sentence/review
              for word in sent: # for each word in a review/sentence
                  if word in w2v_words and word in tfidf_feat1:
                      vec = w2v_model.wv[word]
                      #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
                      # to reduce the computation we are
                      # dictionary[word] = idf value of word in whole corpus
                      # sent.count(word) = tf values of word in this review
                      tf_idf = dictionary1[word]*(sent.count(word)/len(sent))
                      sent_vec += (vec * tf_idf)
                      weight_sum += 1
              if weight_sum != 0:
                  sent_vec /= weight_sum
              train_tfidf_sent_vectors.append(sent_vec)

In [107]: # TF-IDF weighted Word2Vec
          # Vectorizing test dataset.

          test_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in
```

```

for sent in X_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat2:
            vec = w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf values of word in this review
            tf_idf = dictionary2[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum +=1
    if weight_sum != 0:
        sent_vec /= weight_sum
    test_tfidf_sent_vectors.append(sent_vec)

```

```

In [108]: # Calculating training error .
# Because of large amount of data. we are processing data into three batches here.
# After processing all the results of these batches are merged into one .

train_auc_list1 = [] # This contains area under curve value of first batch correspond
train_auc_list2 = [] # for second batch
train_auc_list3 = [] # for third batch

# Testing whole training data at once takes a lot of memory which takes a lot of time
# Therefore we are dividing training data into 3 parts and then we will calculate AUC

x_train_1 = train_tfidf_sent_vectors[0:20000][:] # Row 0 to 19999 and all columns
x_train_2 = train_tfidf_sent_vectors[20000:40000][:]
x_train_3 = train_tfidf_sent_vectors[40000:61441][:]

y_train_1 = Y_train[0:20000][:] # Row 0 to 19999 and all columns
y_train_2 = Y_train[20000:40000][:]
y_train_3 = Y_train[40000:61441][:]

# Calculating training error for first batch
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=30,n_jobs=-1)
    clf.fit(x_train_1,y_train_1)

    pre_probab = clf.predict_proba(x_train_1)[:,-1] # Returns probability of positive class

    auc = roc_auc_score(y_train_1,pre_probab)
    train_auc_list1.append(auc)

# Calculating training error for second batch
for k in tqdm(range(1,30,2)):

```

```
clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=30,n_jobs=
clf.fit(x_train_2,y_train_2)
```

```
pre_probab = clf.predict_proba(x_train_2)[:,1]
```

```
auc = roc_auc_score(y_train_2,pre_probab)
train_auc_list2.append(auc)
```

```
# Calculating training error for third batch
```

```
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=30,n_jobs=
    clf.fit(x_train_3,y_train_3)
```

```
pre_probab = clf.predict_proba(x_train_3)[:,1]
```

```
auc = roc_auc_score(y_train_3,pre_probab)
train_auc_list3.append(auc)
```

```
# Combining results together.
```

```
train_auc_list = [(x+y+z)/3 for x,y,z in zip(train_auc_list1,train_auc_list2,train_auc_list3)]
```

```
0%|
7%|
13%|
20%|
27%|
33%|
40%|
47%|
53%|
60%|
67%|
73%|
80%|
87%|
93%|
100%| 15/15 [06:35<00:00, 26.48s/it]
```

```
0%|
7%|
13%|
20%|
27%|
33%|
40%|
47%|
```

```

53%|                                     | 8/15 [03:12<02:43, 23.31s/it]
60%|                                     | 9/15 [03:35<02:18, 23.07s/it]
67%|                                     | 10/15 [04:05<02:05, 25.19s/it]
73%|                                     | 11/15 [04:28<01:38, 24.69s/it]
80%|                                     | 12/15 [04:51<01:12, 24.14s/it]
87%|                                     | 13/15 [05:12<00:46, 23.20s/it]
93%|      | 14/15 [05:34<00:22, 22.71s/it]
100%|| 15/15 [05:58<00:00, 23.11s/it]

0%|
7%|                                     | 1/15 [00:24<00:00, 23.07s/it]
13%|                                     | 2/15 [00:48<00:00, 23.07s/it]
20%|                                     | 3/15 [01:12<04:49, 24.69s/it]
27%|                                     | 4/15 [01:39<04:35, 25.05s/it]
33%|                                     | 5/15 [02:15<04:42, 28.21s/it]
40%|                                     | 6/15 [02:41<04:08, 27.65s/it]
47%|                                     | 7/15 [03:05<03:32, 26.55s/it]
53%|                                     | 8/15 [03:30<03:02, 26.00s/it]
60%|                                     | 9/15 [03:54<02:32, 25.47s/it]
67%|                                     | 10/15 [04:18<02:04, 25.00s/it]
73%|                                     | 11/15 [04:43<01:40, 25.19s/it]
80%|                                     | 12/15 [05:07<01:14, 24.84s/it]
87%|      | 13/15 [05:31<00:49, 24.57s/it]
93%|      | 14/15 [05:57<00:24, 24.78s/it]
100%|| 15/15 [06:24<00:00, 25.64s/it]

```

In [109]: # 10 fold cross validation using time series splitting

```

auc_list = []

for k in tqdm(range(1,30,2)):
    # KNN Classifier
    clf = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree',leaf_size=30,n_jobs=-1)
    i=0
    auc=0.0
    for train_index,test_index in tscv.split(train_tfidf_sent_vectors):
        x_train = train_sent_vectors[0:train_index[-1]][:] # row 0 to train_index(excluding)
        y_train = Y_train[0:train_index[-1]][:] # row 0 to train_index(excluding)
        x_test = train_sent_vectors[train_index[-1]:test_index[-1]][:] # row from train_index to test_index
        y_test = Y_train[train_index[-1]:test_index[-1]][:] # row from train_index to test_index

        clf.fit(x_train,y_train)

        predict_probab = clf.predict_proba(x_test)[:,-1] # Returns probability of for each class
        i += 1
        auc += roc_auc_score(y_test,predict_probab)

```

```
auc_list.append(auc/i) # Averaging auc for all 10 folds .
```

```


0%|
7%|
13%|
20%|
27%|
33%|
40%|
47%|
53%|
60%|
67%|
73%|
80%|
87%|
93%|
100%| 15/15 [44:01<00:00, 176.22s/it]

```

```

In [110]: # Plotting graph of auc and parameter for training and cross validation error
plot_knn_vs_auc(train_auc_list, auc_list)

```



```

In [111]: 30# Training final model on best auc and taking k = 10

```

```

# Training one model with all the data hangs the PC .
# Therefore we will divide data into 3 parts and then train three separate models.
x_train_1 = train_tfidf_sent_vectors[0:20000][:] # Row 0 to 19999 and all columns
x_train_2 = train_tfidf_sent_vectors[20000:40000][:]
x_train_3 = train_tfidf_sent_vectors[40000:61441][:]

y_train_1 = Y_train[0:20000][:] # Row 0 to 19999 and all columns
y_train_2 = Y_train[20000:40000][:]
y_train_3 = Y_train[40000:61441][:]

final_clf1 = KNeighborsClassifier(n_neighbors=30,algorithm='kd_tree',leaf_size=40,n_
final_clf1.fit(x_train_1,y_train_1)
predict_probab_1 = final_clf1.predict_proba(test_tfidf_sent_vectors)[: ,1] # This ret
predict_y_train1 = list(final_clf1.predict(x_train_1))

final_clf2 = KNeighborsClassifier(n_neighbors=30,algorithm='kd_tree',leaf_size=40,n_
final_clf2.fit(x_train_2,y_train_2)
predict_probab_2 = final_clf2.predict_proba(test_tfidf_sent_vectors)[: ,1] # This ret
predict_y_train2 = list(final_clf1.predict(x_train_2))

final_clf3 = KNeighborsClassifier(n_neighbors=30,algorithm='kd_tree',leaf_size=40,n_
final_clf3.fit(x_train_3,y_train_3)
predict_probab_3 = final_clf3.predict_proba(test_tfidf_sent_vectors)[: ,1] # This ret
predict_y_train3 = final_clf1.predict(x_train_3)

# Now merging all the three probability scores into one

predict_probab = [(x+y+z)/3 for x,y,z in zip(predict_probab_1,predict_probab_2,predi
predict_y = [1 if(x+y+z>=2) else 0 for x,y,z in zip(predict_probab_1,predict_probab_2

# Combining results for train dataset
predict_y_train = np.concatenate((predict_y_train1,predict_y_train2),axis=None)
predict_y_train = np.concatenate((predict_y_train,predict_y_train3),axis=None)

auc = roc_auc_score(Y_test,predict_probab)
print("Final AUC is ::{:.2f}".format(auc))

```

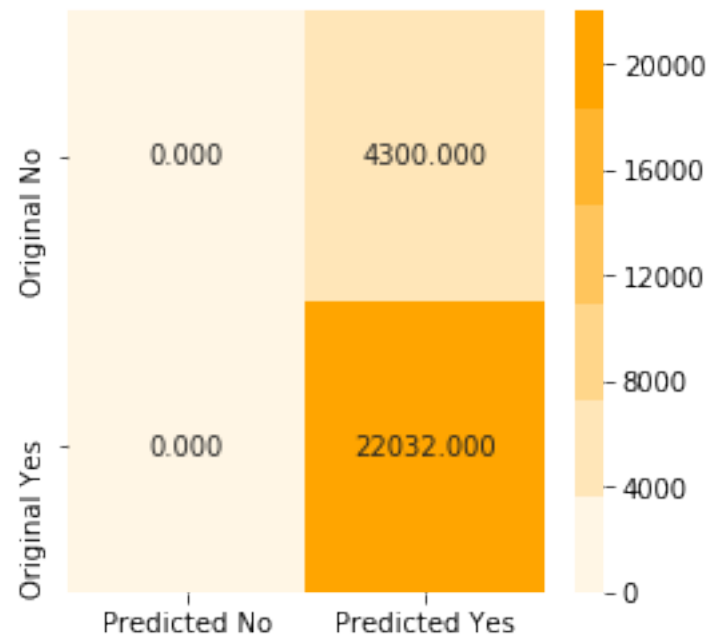
Final AUC is ::0.50

```


In [112]: # Plotting confusion matrix
          confusion_matrix_plot(Y_test,predict_y)

```

Confusion matrix



```
In [113]: # Plotting roc curve
plot_roc_curve(Y_test,predict_y,Y_train,predict_y_train)
```



6.2 [5.2] Applying KNN kd-tree

6.2.1 [5.2.1] Applying KNN kd-tree on BOW, SET 5

```
In [61]: # Kd tree works slow for high dimensional data.
# Therefore taking top 5000 features
from sklearn.cross_validation import train_test_split
count_vect = CountVectorizer(max_features=5000)

X = cleaned_reviews
Y = final["Score"]
# Splitting data into train and test dataset
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.3,random_state=42)
print(len(X_train),len(X_test))
```

61441 26332

```
In [62]: # Vectorizing train and test dataset seperately
X_train = count_vect.fit_transform(X_train)
print(X_train.shape)

X_test = count_vect.transform(X_test)
print(X_test.shape)
```

(61441, 5000)

(26332, 5000)

```
In [63]: # Since kd Tree dont take sparse matrix .
# Therefore we are converting sparse to dense matrix using TruncatedSVD
from sklearn.decomposition import TruncatedSVD

# Initializing TruncatedSVD
# Too many features takes a lot of time . Therefore taking 500 features only
svd = TruncatedSVD(n_components=200,algorithm='randomized',n_iter=50,random_state=42)
X_train = svd.fit_transform(X_train)
X_test = svd.fit_transform(X_test)
print(X_train.shape,X_test.shape)
```

(61441, 200) (26332, 200)

```
In [66]: # Calculating training error .
# Because of large amount of data. we are processing data into three batches here.
# After processing all the results of these batches are merged into one .

param_list = [1,3,5,7,9,11,13,15,17,19,21,23,25,27,29]
train_auc_list1 = [] # This contains area under curve value of first batch correspon
train_auc_list2 = [] # for stqdm(econd batch
```

```

train_auc_list3 = []    # for third batch

# Testing whole training data at once takes a lot of memory which takes a lot of time
# There fore we are dividing training data into 3 parts and then we will calculate AU

x_train_1 = X_train[0:20000][:] # Row 0 to 19999 and all columns
x_train_2 = X_train[20000:40000][:]
x_train_3 = X_train[40000:61441][:]

y_train_1 = Y_train[0:20000][:] # Row 0 to 19999 and all columns
y_train_2 = Y_train[20000:40000][:]
y_train_3 = Y_train[40000:61441][:]

# Calculating training error for first batch
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=40)
    clf.fit(x_train_1,y_train_1)

    pre_probab = clf.predict_proba(x_train_1)[:,:1] # Returns probability of positive

    auc = roc_auc_score(y_train_1,pre_probab)
    train_auc_list1.append(auc)

# Calculating training error for second batch
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=40)
    clf.fit(x_train_2,y_train_2)

    pre_probab = clf.predict_proba(x_train_2)[:,:1]

    auc = roc_auc_score(y_train_2,pre_probab)
    train_auc_list2.append(auc)

# Calculating training error for third batch
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=40)
    clf.fit(x_train_3,y_train_3)

    pre_probab = clf.predict_proba(x_train_3)[:,:1]

    auc = roc_auc_score(y_train_3,pre_probab)
    train_auc_list3.append(auc)

100%|| 15/15 [2:16:40<00:00, 595.11s/it]
100%|| 15/15 [1:52:25<00:00, 427.10s/it]
100%|| 15/15 [1:02:25<00:00, 273.88s/it]

```

```

In [67]: train_auc_list = [(x+y+z)/3 for x,y,z in zip(train_auc_list1,train_auc_list2,train_auc_list3)]

In [68]: # Performing 10 fold cross validation on time split data

tscv = TimeSeriesSplit(n_splits=10)
auc_list1 = []

for k in range(1,30,2):
    clf1 = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree',leaf_size=40)
    auc1 = 0.0
    i1=0
    for train_index,test_index in tscv.split(X_train):
        x_train = X_train[0:train_index[-1]][:] # row 0 to train_index(excluding)
        y_train = Y_train[0:train_index[-1]][:] # row 0 to train_index(excluding)
        x_test = X_train[train_index[-1]:test_index[-1]][:] # row from train_index to test_index(excluding)
        y_test = Y_train[train_index[-1]:test_index[-1]][:] # row from train_index to test_index(excluding)

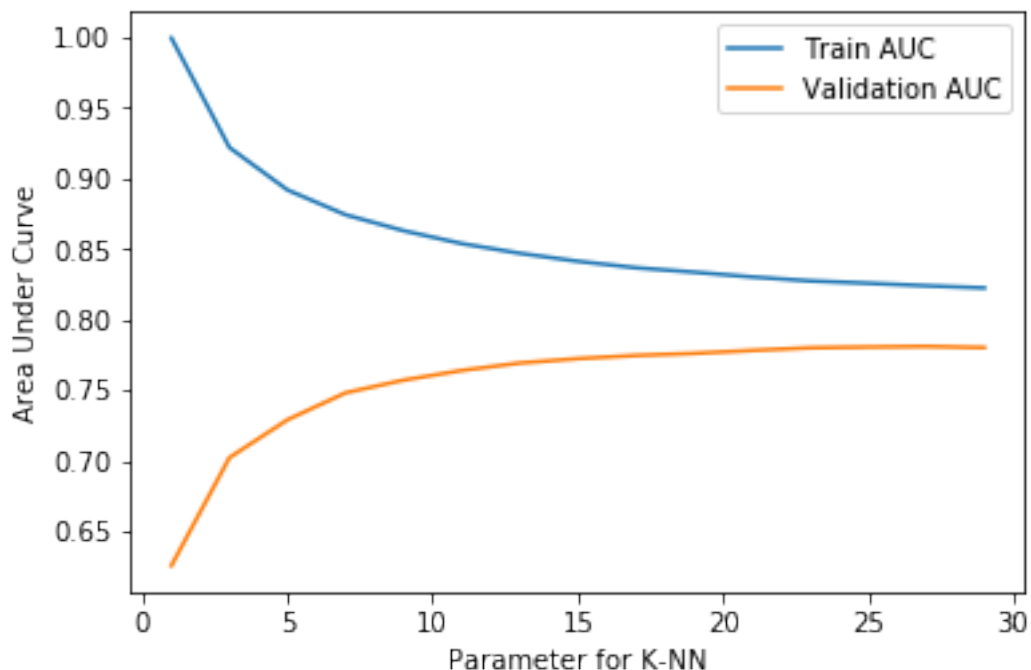
        clf1.fit(x_train,y_train)

        predict_probab1 = clf1.predict_proba(x_test)[:,-1]
        i1 += 1
        auc1 += roc_auc_score(y_test,predict_probab1)

    auc_list1.append(auc1/i1)

In [70]: # Plotting graph of auc and parameter for training and cross validation error
plot_knn_vs_auc(train_auc_list,auc_list1)

```



```
In [75]: # Training final model on best auc and taking k = 25
```

```
final_clf = KNeighborsClassifier(n_neighbors=25,algorithm='kd_tree',leaf_size=40)
final_clf.fit(X_train,Y_train)
predict_probab = final_clf.predict_proba(X_test)[:,-1] # This returns only probability
predict_y = final_clf.predict(X_test)


auc = roc_auc_score(Y_test,predict_probab)
print("Final AUC is ::{:.2f}".format(auc))
```

Final AUC is ::0.63


```
In [72]: predict_y_train = final_clf.predict(X_train)
```

```
In [76]: # Plotting confusion matrix
confusion_matrix_plot(Y_test,predict_y)
```

Confusion matrix



```
In [77]: # Plotting roc curve
plot_roc_curve(Y_test,predict_y,Y_train,predict_y_train)
```



6.2.2 [5.2.2] Applying KNN kd-tree on TFIDF, SET 6

In [78]: *# In this section Tfidf will be used for vectorization*
Splitting datasets into train and test datasets

```
X = final['Cleaned_review']
Y = final['Score']
tf_idf_vect = TfidfVectorizer(max_features=300)
```

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.3,random_state=42)
```

```
# Now we will vectorize train and test datasets separately using Tfidf
# Use fit_transform to vectorize train dataset and transform to vectorize test dataset
X_train = tf_idf_vect.fit_transform(X_train)
X_test = tf_idf_vect.transform(X_test)
```

```
In [79]: from sklearn.decomposition import TruncatedSVD
svd = TruncatedSVD(n_components=200,algorithm='randomized',n_iter=50,random_state=42)
X_train = svd.fit_transform(X_train)
X_test = svd.fit_transform(X_test)
print(X_train.shape,X_test.shape)
```

```
(61441, 200) (26332, 200)
```

```

In [80]: param_list = [1,3,5,7,9,11,13,15,17,19,21,23,25,27,29]
        train_auc_list1 = []    # This contains area under curve value of first batch correspond
        train_auc_list2 = []    # for second batch
        train_auc_list3 = []    # for third batch

        # Testing whole training data at once takes a lot of memory which takes a lot of time
        # There fore we are dividing training data into 3 parts and then we will calculate AU

        x_train_1 = X_train[0:20000][:] # Row 0 to 19999 and all columns
        x_train_2 = X_train[20000:40000][:]
        x_train_3 = X_train[40000:61441][:]

        y_train_1 = Y_train[0:20000][:] # Row 0 to 19999 and all columns
        y_train_2 = Y_train[20000:40000][:]
        y_train_3 = Y_train[40000:61441][:]

        # Calculating training error for first batch
        for k in tqdm(range(1,30,2)):
            clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=40,n_jobs=-1)
            clf.fit(x_train_1,y_train_1)

            pre_probab = clf.predict_proba(x_train_1)[:,-1] # Returns probability of positive

            auc = roc_auc_score(y_train_1,pre_probab)
            train_auc_list1.append(auc)

        # Calculating training error for second batch
        for k in tqdm(range(1,30,2)):
            clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=40,n_jobs=-1)
            clf.fit(x_train_2,y_train_2)

            pre_probab = clf.predict_proba(x_train_2)[:,-1]

            auc = roc_auc_score(y_train_2,pre_probab)
            train_auc_list2.append(auc)

        # Calculating training error for third batch
        for k in tqdm(range(1,30,2)):
            clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=40,n_jobs=-1)
            clf.fit(x_train_3,y_train_3)

            pre_probab = clf.predict_proba(x_train_3)[:,-1]

            auc = roc_auc_score(y_train_3,pre_probab)
            train_auc_list3.append(auc)

```

100%|| 15/15 [19:36<00:00, 84.13s/it]

```
100%| 15/15 [20:11<00:00, 90.45s/it]
100%| 15/15 [27:46<00:00, 117.68s/it]
```

```
In [81]: train_auc_list = [(x+y+z)/3 for x,y,z in zip(train_auc_list1,train_auc_list2,train_auc_list3)]
```

```
In [84]: # Performing 10 fold cross validation on time split data
```

```
tscv = TimeSeriesSplit(n_splits=10)
auc_list1 = []

for k in tqdm(range(1,30,2)):
    clf1 = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree',leaf_size=40,n_jobs=-1)
    auc1 = 0.0
    i1=0
    for train_index,test_index in tscv.split(X_train):
        x_train = X_train[0:train_index[-1]][:] # row 0 to train_index(excluding)
        y_train = Y_train[0:train_index[-1]][:] # row 0 to train_index(excluding)
        x_test = X_train[train_index[-1]:test_index[-1]][:] # row from train_index to test_index(excluding)
        y_test = Y_train[train_index[-1]:test_index[-1]][:] # row from train_index to test_index(excluding)

        clf1.fit(x_train,y_train)

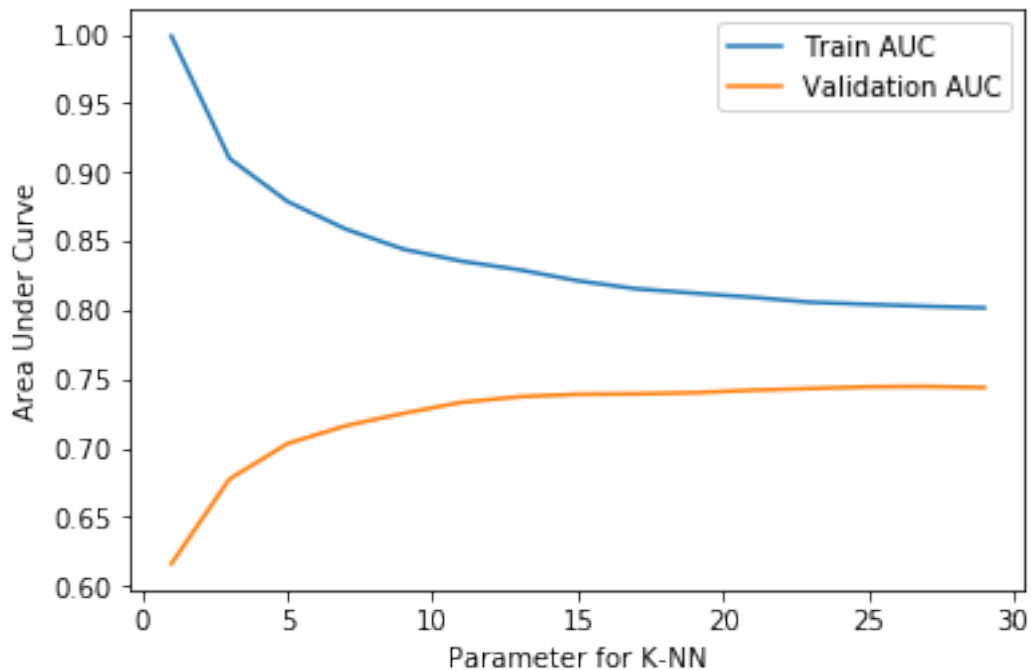
        predict_probab1 = clf1.predict_proba(x_test)[:,-1]
        i1 += 1
        auc1 += roc_auc_score(y_test,predict_probab1)

    auc_list1.append(auc1/i1)
```

```
0%|
7%|
13%|
20%|
27%|
33%|
40%|
47%|
53%|
60%|
67%|
73%|
80%|
87%|
93%|
100%| 15/15 [1:52:53<00:00, 455.96s/it]
```



```
In [85]: # Plotting graph of auc and parameter for training and cross validation error
plot_knn_vs_auc(train_auc_list,auc_list1)
```



```
In [86]: # Training final model on best auc and taking k = 11
```

```
final_clf = KNeighborsClassifier(n_neighbors=24,algorithm='kd_tree',leaf_size=30,n_jobs=4)

final_clf.fit(X_train,Y_train)


predict_probab = final_clf.predict_proba(X_test)[:,-1] # This returns only probability
predict_y = final_clf.predict(X_test)
auc = roc_auc_score(Y_test,predict_probab)
print("Final AUC is {:.2f}".format(auc))
```

Final AUC is 0.57


```
In [89]: predict_y_train = final_clf.predict(X_train)
```

```
In [87]: # Plotting confusion matrix
confusion_matrix_plot(Y_test,predict_y)
```

Confusion matrix



```
In [90]: # Plotting roc curve
plot_roc_curve(Y_test,predict_y,Y_train,predict_y_train)
```



6.2.3 [5.2.3] Applying KNN kd-tree on AVG W2V, SET 3

```
In [91]: # Splitting data into train and test.
        Y = final['Score']
        X_train,X_test,Y_train,Y_test = train_test_split(list_of_sentence,Y,test_size=0.3,ran

In [92]: # Now we will vectorize train dataset usin avg_w2v
        train_sent_vectors = []; # the avg-w2v for each sentence/review is stored in this lis
        for sent in X_train: # for each review/sentence
            sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need t
            cnt_words =0; # num of words with a valid vector in the sentence/review
            for word in sent: # for each word in a review/sentence
                if word in w2v_words:
                    vec = w2v_model.wv[word]
                    sent_vec += vec
                    cnt_words += 1
            if cnt_words != 0:
                sent_vec /= cnt_words
            train_sent_vectors.append(sent_vec)
        print(len(train_sent_vectors))
        print(len(train_sent_vectors[0]))
```

61441

50

```
In [93]: # Vectorization of test dataset using avg_w2v

        test_sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
        for sent in X_test: # for each review/sentence
            sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need t
            cnt_words =0; # num of words with a valid vector in the sentence/review
            for word in sent: # for each word in a review/sentence
                if word in w2v_words:
                    vec = w2v_model.wv[word]
                    sent_vec += vec
                    cnt_words += 1
            if cnt_words != 0:
                sent_vec /= cnt_words
            test_sent_vectors.append(sent_vec)
        print(len(test_sent_vectors))
        print(len(test_sent_vectors[0]))
```

26332

50

```
In [94]: param_list = [1,3,5,7,9,11,13,15,17,19,21,23,25,27,29]
        train_auc_list1 = [] # This contains area under curve value of first batch correspon
```

```

train_auc_list2 = []    # for second batch
train_auc_list3 = []    # for third batch

# Testing whole training data at once takes a lot of memory which takes a lot of time
# There fore we are dividing training data into 3 parts and then we will calculate AU

x_train_1 = train_sent_vectors[0:20000][:] # Row 0 to 19999 and all columns
x_train_2 = train_sent_vectors[20000:40000][:]
x_train_3 = train_sent_vectors[40000:61441][:]

y_train_1 = Y_train[0:20000][:] # Row 0 to 19999 and all columns
y_train_2 = Y_train[20000:40000][:]
y_train_3 = Y_train[40000:61441][:]

# Calculating training error for first batch
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=30,n_jobs=-1)
    clf.fit(x_train_1,y_train_1)

    pre_probab = clf.predict_proba(x_train_1)[:,-1] # Returns probability of positive class

    auc = roc_auc_score(y_train_1,pre_probab)
    train_auc_list1.append(auc)

# Calculating training error for second batch
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=30,n_jobs=-1)
    clf.fit(x_train_2,y_train_2)

    pre_probab = clf.predict_proba(x_train_2)[:,-1]

    auc = roc_auc_score(y_train_2,pre_probab)
    train_auc_list2.append(auc)

# Calculating training error for third batch
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=30,n_jobs=-1)
    clf.fit(x_train_3,y_train_3)

    pre_probab = clf.predict_proba(x_train_3)[:,-1]

    auc = roc_auc_score(y_train_3,pre_probab)
    train_auc_list3.append(auc)

# Combining the result into one.
train_auc_list = [(x+y+z)/3 for x,y,z in zip(train_auc_list1,train_auc_list2,train_auc_list3)]

```

0%|
 7%| | 1/15 [00:00:00, 0.00s/it]
 13%| | 2/15 [00:39<02:00, 0.26s/it]
 20%| | 3/15 [01:19<04:03, 20.27s/it]
 27%| | 4/15 [01:59<04:50, 26.41s/it]
 33%| | 5/15 [02:39<05:05, 30.53s/it]
 40%| | 6/15 [03:19<04:59, 33.30s/it]
 47%| | 7/15 [04:00<04:43, 35.42s/it]
 53%| | 8/15 [04:42<04:22, 37.52s/it]
 60%| | 9/15 [05:23<03:51, 38.58s/it]
 67%| | 10/15 [06:05<03:17, 39.48s/it]
 73%| | 11/15 [06:46<02:40, 40.17s/it]
 80%| | 12/15 [07:28<02:01, 40.64s/it]
 87%| | 13/15 [08:09<01:21, 40.60s/it]
 93%| | 14/15 [08:59<00:43, 43.47s/it]
 100%| | 15/15 [09:50<00:00, 45.65s/it]

0%|
 7%| | 1/15 [00:00:00, 0.00s/it]
 13%| | 2/15 [00:47<03:00, 0.31s/it]
 20%| | 3/15 [01:36<04:57, 24.47s/it]
 27%| | 4/15 [02:25<05:52, 32.04s/it]
 33%| | 5/15 [03:14<06:10, 37.08s/it]
 40%| | 6/15 [04:03<06:07, 40.79s/it]
 47%| | 7/15 [04:53<05:47, 43.50s/it]
 53%| | 8/15 [05:42<05:16, 45.19s/it]
 60%| | 9/15 [06:32<04:39, 46.54s/it]
 67%| | 10/15 [07:23<03:59, 47.92s/it]
 73%| | 11/15 [08:15<03:16, 49.11s/it]
 80%| | 12/15 [09:05<02:27, 49.32s/it]
 87%| | 13/15 [09:56<01:39, 49.92s/it]
 93%| | 14/15 [10:46<00:49, 49.84s/it]
 100%| | 15/15 [11:36<00:00, 50.08s/it]

0%|
 7%| | 1/15 [00:00:00, 0.00s/it]
 13%| | 2/15 [00:53<03:00, 0.35s/it]
 20%| | 3/15 [01:46<05:28, 27.47s/it]
 27%| | 4/15 [02:42<06:33, 35.79s/it]
 33%| | 5/15 [03:38<06:58, 41.90s/it]
 40%| | 6/15 [04:35<06:59, 46.59s/it]
 47%| | 7/15 [05:31<06:33, 49.21s/it]
 53%| | 8/15 [06:28<06:02, 51.78s/it]
 60%| | 9/15 [07:26<05:20, 53.40s/it]
 67%| | 10/15 [08:23<04:33, 54.67s/it]
 73%| | 11/15 [09:17<03:37, 54.47s/it]
 80%| | 12/15 [10:12<02:43, 54.45s/it]

```

87%|          | 13/15 [11:08<01:50, 55.12s/it]
93%|          | 14/15 [12:06<00:55, 55.89s/it]
100%|| 15/15 [13:05<00:00, 56.76s/it]

```

```


In [95]: # 10 fold cross validation using time series splitting
# Here X_train is train_sent_vectors and X_test is test_sent_vectors
# Here vectorization results of previous section is used . Only KNN algorithm is chan
tscv = TimeSeriesSplit(n_splits=10)
auc_list=[]
for k in range(1,30,2):
    # KNN Classifier
    clf = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree',leaf_size=40,n_jobs=
    i=0
    auc=0.0
    for train_index,test_index in tscv.split(train_sent_vectors):
        x_train = train_sent_vectors[0:train_index[-1]][:] # row 0 to train_index(exc
        y_train = Y_train[0:train_index[-1]][:] # row 0 to train_index(excluding)
        x_test = train_sent_vectors[train_index[-1]:test_index[-1]][:] # row from tra
        y_test = Y_train[train_index[-1]:test_index[-1]][:] # row from train_index to

        clf.fit(x_train,y_train)
        predict_probab = clf.predict_proba(x_test)[:,-1]
        #print(len(x_test),len(y_test),len(predict_probab))
        i += 1
        auc += roc_auc_score(y_test,predict_probab)

    auc_list.append(auc/i)

In [96]: # Plotting graph of auc and parameter for training and cross validation error
plot_knn_vs_auc(train_auc_list, auc_list)

```



```
In [97]: # Training the final model
final_clf = KNeighborsClassifier(n_neighbors=25,algorithm='kd_tree',leaf_size=30,n_jobs=-1)

final_clf.fit(train_sent_vectors,Y_train)


predict_probab = final_clf.predict_proba(test_sent_vectors)[:,-1] # Returns the class probabilities
predict_y = final_clf.predict(test_sent_vectors)
predict_y_train = final_clf.predict(train_sent_vectors)
auc = roc_auc_score(Y_test,predict_probab)

print("Final AUC is {:.2f}".format(auc))
```

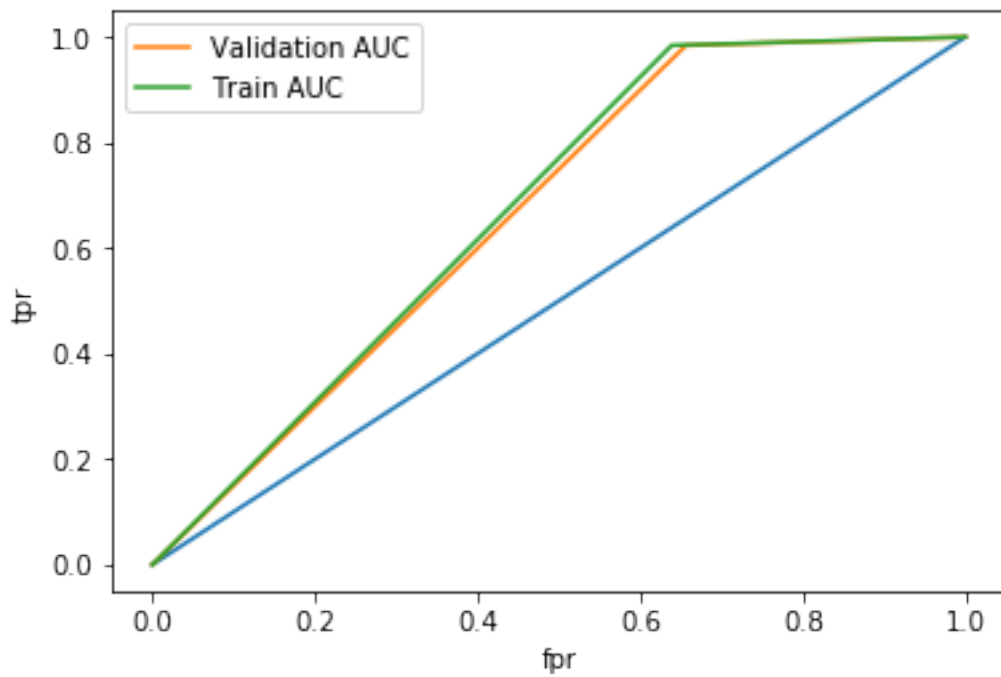
Final AUC is 0.89

```
In [98]: # Plotting confusion matrix
confusion_matrix_plot(Y_test,predict_y)
```

Confusion matrix



```
In [101]: ## Plotting roc curve
plot_roc_curve(Y_test,predict_y,Y_train,predict_y_train)
```



6.2.4 [5.2.4] Applying KNN kd-tree on TFIDF W2V, SET 4

In [114]: *# Using vectorized data from previous section*

```
# Calculating training error .
# Because of large amount of data. we are processing data into three batches here.
# After processing all the results of these batches are merged into one .

train_auc_list1 = [] # This contains area under curve value of first batch corresp
train_auc_list2 = [] # for second batch
train_auc_list3 = [] # for third batch

# Testing whole training data at once takes a lot of memory which takes a lot of time
# There fore we are dividing training data into 3 parts and then we will calculate A

x_train_1 = train_tfidf_sent_vectors[0:20000][:] # Row 0 to 19999 and all columns
x_train_2 = train_tfidf_sent_vectors[20000:40000][:]
x_train_3 = train_tfidf_sent_vectors[40000:61441][:]

y_train_1 = Y_train[0:20000][:] # Row 0 to 19999 and all columns
y_train_2 = Y_train[20000:40000][:]
y_train_3 = Y_train[40000:61441][:]

# Calculating training error for first batch
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=40,n_jobs=
    clf.fit(x_train_1,y_train_1)

    pre_probab = clf.predict_proba(x_train_1)[:,:1] # Returns probability of positive

    auc = roc_auc_score(y_train_1,pre_probab)
    train_auc_list1.append(auc)

# Calculating training error for second batch
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=40,n_jobs=
    clf.fit(x_train_2,y_train_2)

    pre_probab = clf.predict_proba(x_train_2)[:,:1]

    auc = roc_auc_score(y_train_2,pre_probab)
    train_auc_list2.append(auc)

# Calculating training error for third batch
for k in tqdm(range(1,30,2)):
    clf = KNeighborsClassifier(n_neighbors=k,algorithm="kd_tree",leaf_size=40,n_jobs=
    clf.fit(x_train_3,y_train_3)
```

```

pre_probab = clf.predict_proba(x_train_3)[: ,1]

auc = roc_auc_score(y_train_3,pre_probab)
train_auc_list3.append(auc)

# Combining results together.
train_auc_list = [(x+y+z)/3 for x,y,z in zip(train_auc_list1,train_auc_list2,train_auc_list3)]

```

```

0%|
7%|
13%|
20%|
27%|
33%|
40%|
47%|
53%|
60%|
67%|
73%|
80%|
87%|
93%|
100%| 15/15 [05:10<00:00, 21.46s/it]

```

```

0%|
7%|
13%|
20%|
27%|
33%|
40%|
47%|
53%|
60%|
67%|
73%|
80%|
87%|
93%|
100%| 15/15 [05:21<00:00, 21.28s/it]

```

```

0%|
7%|
13%|
20%|

```

| | | |
|------|--------------------------------|--------------------------------|
| 27% | | 4/15 [01:49<05:03, 27.59s/it] |
| 33% | | 5/15 [02:17<04:35, 27.55s/it] |
| 40% | | 6/15 [02:43<04:05, 27.26s/it] |
| 47% | | 7/15 [03:11<03:38, 27.27s/it] |
| 53% | | 8/15 [03:39<03:12, 27.51s/it] |
| 60% | | 9/15 [04:06<02:45, 27.52s/it] |
| 67% | | 10/15 [04:34<02:17, 27.50s/it] |
| 73% | | 11/15 [05:01<01:49, 27.27s/it] |
| 80% | | 12/15 [05:27<01:21, 27.07s/it] |
| 87% | | 13/15 [05:54<00:54, 27.08s/it] |
| 93% | 14/15 [06:21<00:27, 27.05s/it] | |
| 100% | 15/15 [06:50<00:00, 27.43s/it] | |

```

In [115]: # 10 fold cross validation using time series splitting
from sklearn.model_selection import TimeSeriesSplit
tscv = TimeSeriesSplit()
auc_list = []

for k in range(1,30,2):
    # KNN Classifier
    clf = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree',leaf_size=40,n_jobs=
    i=0
    auc=0.0
    for train_index,test_index in tscv.split(train_tfidf_sent_vectors):
        x_train = train_tfidf_sent_vectors[0:train_index[-1]][:] # row 0 to train_index
        y_train = Y_train[0:train_index[-1]][:] # row 0 to train_index(excluding)
        x_test = train_tfidf_sent_vectors[train_index[-1]:test_index[-1]][:] # row from
        y_test = Y_train[train_index[-1]:test_index[-1]][:] # row from train_index to

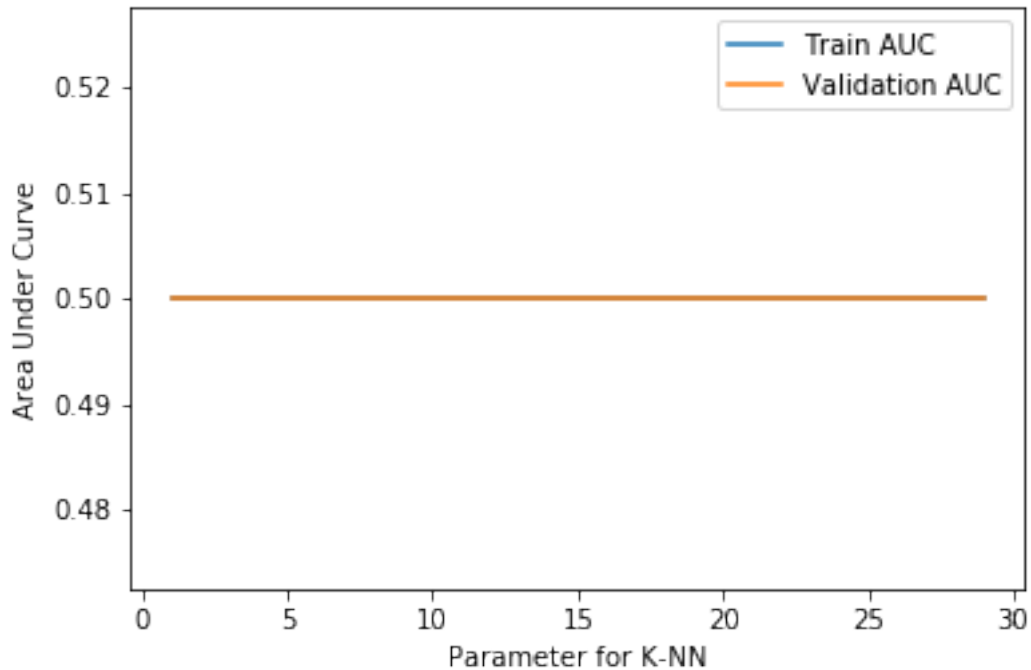
        clf.fit(x_train,y_train)

        predict_probab = clf.predict_proba(x_test)[:,-1] # Returns probability of for
        i += 1
        auc += roc_auc_score(y_test,predict_probab)

    auc_list.append(auc/i) # Averaging auc for all 10 folds .

In [117]: # Plotting graph of auc and parameter for training and cross validation error
plot_knn_vs_auc(train_auc_list,auc_list)

```



```
In [118]: # Training the final model
          final_clf = KNeighborsClassifier(n_neighbors=11,algorithm='kd_tree',leaf_size=40,n_j
          final_clf.fit(train_tfidf_sent_vectors,Y_train)

          predict_probab = final_clf.predict_proba(test_tfidf_sent_vectors)[: ,1] # Returns the
          auc = roc_auc_score(Y_test,predict_probab)


          print("AUC of model is {:.2f}".format(auc))
```

AUC of model is 0.50


```
In [119]: # Train and test prediction
          predict_y = final_clf.predict(test_tfidf_sent_vectors)
          predict_y_train = final_clf.predict(train_tfidf_sent_vectors)
```

```
In [120]: # Plotting confusion matrix , precision and recall matrix
          confusion_matrix_plot(Y_test,predict_y)
```

Confusion matrix



```
In [121]: # Plotting ROC Curve
plot_roc_curve(Y_test,predict_y,Y_train,predict_y_train)
```



7 [6] Conclusions

```
In [1]: from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Model Type", "Best K", "AUC",]

x.add_row(["BOW", "21", "0.67"])
x.add_row(["TfIdf", "25", "0.51"])
x.add_row(["Avg W2V", "17", "0.87"])
x.add_row(["TfIdf Weighted W2V", "21", "0.50"])
print(x)
```

```
+-----+-----+-----+
|   Model Type   | Best K | AUC  |
+-----+-----+-----+
|      BOW      |    21  | 0.67 |
|     TfIdf     |    25  | 0.51 |
|    Avg W2V    |    17  | 0.87 |
| TfIdf Weighted W2V |    21  | 0.50 |
+-----+-----+-----+
```